Open in app ↗ Sign up Sign in

Medium Search Write

The Ultimate AWS Certified Machine Learning Engineer — Associate MLA-C01 cheatsheet: Tools, Techniques, and Key Concepts

Rafa Xu Follow 7 min read · Jan 26, 2025

1 1

Are you looking for an AWS Certified Machine Learning Engineer — Associate MLA-C01 cheatsheet*? This cheatsheet consolidates various AWS tools and ML fundamentals into convenient categories, so you can quickly find what you need Let's dive in!*

1. The Amazon SageMaker Ecosystem

1.1 SageMaker Studio

- **Amazon SageMaker Studio**: An integrated development environment (IDE) for machine learning, providing a single, web-based interface to perform data preprocessing, model training, and deployment.

1.2 SageMaker Domain & User Profile

- **SageMaker Domain**: A domain captures user settings, configurations, and security policies for Studio usage across your organization.
- **SageMaker User Profile**: Represents an individual user within a SageMaker Domain, storing preferences like kernels and compute settings.

1.3 SageMaker Pipelines

- **Amazon SageMaker Pipelines**: A fully managed workflow service for building, automating, and managing ML pipelines — handling tasks like data loading, model training, and deployment in a CI/CD-like fashion.

1.4 SageMaker Autopilot

- **SageMaker Autopilot**: Automatically explores multiple ML algorithms, tuning parameters to create high-quality models. Provides generated notebooks for transparency and reproducibility.

1.5 SageMaker Model Registry & Model Group

- **Model Registry**: A central repository to track and catalog trained models, maintaining version control, approvals, and deployment history.
- **Model Group**: A logical container grouping together different versions of the same model, facilitating model versioning and promotion to production.

1.6 SageMaker ML Lineage Tracking

- **ML Lineage**: Tracks the origin of datasets, transformations, and training runs that lead to model creation, helping with reproducibility and audit trails.

1.7 SageMaker Managed Spot Training

- **Managed Spot Training**: Leverages unused EC2 capacity at a lower cost. Automatically retries interrupted jobs to reduce overall training costs while handling large datasets.

1.8 SageMaker Distributed Data Parallel (SMDDP)

- **SMDDP Library**: A distributed training library enabling data parallelism across multiple instances, thus reducing training time for large models or datasets.

1.9 SageMaker Experiments

- **Experiments**: A feature that lets you organize, track, compare, and evaluate machine learning experiments (training jobs) at scale within SageMaker.

1.10 SageMaker Debugger

- **Debugger**: Monitors model training in real time, capturing metrics (e.g., gradients, weights) and providing alerts or rules to detect anomalies like vanishing gradients.

1.11 SageMaker Endpoint Production Variants

- **Production Variants:** Enables deploying multiple variants of a model to a single SageMaker endpoint, facilitating A/B testing and canary releases for model updates.

1.12 SageMaker Shadow Testing

- **Shadow Testing:** Allows you to send a copy of real-time inference traffic to a 'shadow' endpoint, helping validate new model versions without affecting production outcomes.

1.13 SageMaker Inference Recommender

- **Inference Recommender:** Helps you choose the optimal instance type and configuration for model inference, automating load testing and performance analysis.

1.14 SageMaker Canvas

- **Canvas:** A no-code/low-code interface for business analysts to build ML models via a visual interface, without requiring extensive coding knowledge.

1.15 SageMaker Script Mode & Pre-built Containers

- **Script Mode:** Lets you bring your own training script while benefiting from SageMaker's high-performance and distributed training environment.
- **Pre-built Containers:** Managed Docker containers with preinstalled frameworks (e.g., TensorFlow, PyTorch, MXNet) for quick training and deployment.

2. Data Preparation & Feature Engineering

2.1 Amazon SageMaker Data Wrangler

- **Data Wrangler**: A tool within SageMaker Studio for data preparation. You can import datasets, clean and transform data, and automate feature engineering flows.

2.2 AWS Glue & Glue DataBrew

- **AWS Glue**: A fully managed ETL service to prepare, transform, and load data for analytics.
- **Glue DataBrew**: A visual data preparation tool for cleaning and normalizing data with an interactive interface.

2.5 Feature Engineering Techniques

- **Feature Group**: In SageMaker Feature Store, a logical grouping of data that includes feature definitions and related metadata.
- **One-Hot Encoding**: Transforms categorical variables into a series of binary (0/1) features.
- **Feature Splitting**: Breaks down composite features (e.g., text fields) into multiple meaningful features (e.g., parsing or tokenizing).
- **Logarithmic Transformation**: Applies log(x)\log(x) to reduce skewness and stabilize variance.
- **Standardized Distribution**: Scales features to have zero mean and unit variance.

- **Padding**: Common in NLP or image tasks to make all input data consistent in size.

2.6 Data Wrangler Image Transformations

- **Enhance Image Contrast**: Adjusts image brightness and contrast for better visibility or model performance.
- **Corrupt Image Transform**: Purposefully degrades image quality to augment data and improve model robustness.
- **Resize Image Transform**: Uniformly changes the resolution of images while preserving aspect ratios.

3. Training & Tuning

3.1 Hyperparameter Tuning

- **Hyperband**: An optimization algorithm that dynamically allocates resources to promising hyperparameter configurations.
- **Grid Search**: Exhaustively tests all combinations of hyperparameters within specified ranges.
- **Random Search**: Samples random hyperparameter combinations, often faster than grid search.
- **Bayesian Optimization**: Builds a probabilistic model of the objective function, intelligently picking hyperparameters based on past evaluation results.

3.2 Key Training Concepts

- **Learning Rate**: Controls how quickly or slowly a model learns; too high can overshoot minima, too low can slow convergence.
- **Early Stopping**: Halts training when validation performance ceases to improve, preventing overfitting.
- **Dropout Layers**: Randomly "drops" neurons during training to reduce overfitting in neural networks.

3.3 SageMaker Managed Warm Pods

- **Managed Warm Pods**: Pre-initialized containers in SageMaker (similar concept to serverless warm starts) that reduce cold-start times for training or inference.

3.4 AWS Trainium & AWS Inferentia

- **AWS Trainium**: Custom ML training chips designed by AWS to deliver high-performance training at lower costs.
- **AWS Inferentia**: Purpose-built chips optimized for high-performance ML inference in the cloud.

3.5 EMR Cluster Nodes

- **Primary, Core, and Task Nodes**: In Amazon EMR, the *primary* node manages the cluster, *core* nodes store data in HDFS and run tasks, while *task* nodes only run tasks without storing data.

4. Inference & Deployment

4.1 SageMaker Real-Time Inference

- **Real-Time Inference**: Deploy machine learning models to an HTTPS endpoint for sub-second latency predictions.

Get Rafa Xu's stories in your inbox

Join Medium for free to get updates from this writer.

Enter your email

Subscribe

Remember me for faster sign in

4.2 SageMaker Batch Transform

- **Batch Transform**: Asynchronously runs high-volume inference on batch data stored in Amazon S3 without needing a persistent endpoint.

4.3 SageMaker Serverless Inference

- **Serverless Inference**: Automatically provisions and scales compute resources to handle inference requests on demand, ideal for low-traffic or unpredictable workloads.

4.4 SageMaker Asynchronous Inference

- **Asynchronous Inference**: Receives inference requests in a queue and processes them asynchronously, suitable for large payloads or long processing times.

5. Monitoring & Debugging

5.1 Model Monitor

- **Model Monitor**: Automatically detects data and model-quality drift in production endpoints. You can configure alerts to identify anomalies early.

5.2 Concept Drift

- **Concept Drift**: Occurs when the statistical properties of the target variable or features change over time, causing model degradation if left unchecked.

5.3 SageMaker Debugger

(Mentioned above, but worth repeating under monitoring for emphasis.)

- **Debugger**: Captures model metrics during training, providing alerts if anomalies like vanishing gradients, overfitting, or poor GPU utilization occur.

6. Large Language Models (LLMs)

6.1 Fundamental LLM Concepts

- **Embedding**: A learned representation of tokens or sequences in a continuous vector space.
- **Tokens**: Units of text (words, subwords, or characters) used by language models during training and inference.

- **Temperature**: Controls randomness in text generation; higher values produce more diverse outputs, lower values produce more deterministic outputs.
- **top_k / top_p (Nucleus Sampling)**: Sampling strategies that limit the model's next-token choices to the top-k predictions or the smallest set of top-p cumulative probability tokens.

6.2 Retrieval Augmented Generation (RAG)

- **RAG**: Combines external knowledge retrieval (e.g., from a database or knowledge base) with a generative model, enhancing contextual accuracy and reducing hallucinations.

7. AWS AI Services

7.1 Amazon Comprehend

- **Comprehend**: A natural language processing (NLP) service to extract insights from text (entity recognition, sentiment analysis, topic modeling, etc.).

7.2 Amazon Transcribe

- **Transcribe**: Automatic speech recognition (ASR) service that converts audio to text, supporting multiple languages and real-time transcription.

7.3 Amazon Translate

- **Translate**: Provides real-time and batch text translation, leveraging neural machine translation models for high-quality multilingual output.

7.4 Amazon Rekognition

- **Rekognition**: Image and video analysis service for object detection, facial recognition, and content moderation.

7.5 Amazon Macie

- **Macie**: A fully managed data security and privacy service, using ML to identify, categorize, and protect sensitive data (e.g., PII) in AWS.

7.6 Amazon Mechanical Turk

- **Mechanical Turk**: A crowdsourcing platform to label and annotate data at scale, helpful for training supervised ML models.

7.7 Additional Services (Recommended)

- **Amazon Kendra**: Intelligent search service powered by ML for enterprise data search.
- **Amazon Forecast**: Uses ML to generate accurate forecasts (e.g., product demand, resource needs).
- **Amazon Personalize**: Provides personalized recommendations, leveraging real-time user data.
- **Amazon Lookout for Vision / Metrics**: Services for anomaly detection in images and numerical data, respectively.

8. Built-In Algorithms & Common ML Tasks

8.1 Amazon SageMaker Built-in Algorithms

- **LightGBM**: Gradient boosting framework for classification and regression tasks, optimized for speed and memory usage.
- **Linear Learner**: Performs linear regression or classification with automatic tuning of model weights.
- **K-Means Clustering**: Unsupervised algorithm for grouping data points into K clusters.
- **Neural Topic Model (NTM)**: Identifies topics in documents using neural network–based topic modeling.
- **Random Cut Forest**: Anomaly detection algorithm for streaming data, also available in Amazon Managed Service for Apache Flink.

8.2 Common ML Tasks

- **Anomaly Detection**: Identifies unusual patterns, outliers, or suspicious events.
- **Linear Regression**: Predicts a continuous value from input features.
- **Logistic Regression**: Estimates the probability of a categorical outcome (e.g., binary classification).
- **Semantic Segmentation**: Assigns labels to each pixel in an image, commonly used in computer vision.
- **Dimensionality Reduction**: Reduces the number of features (e.g., PCA) to mitigate the curse of dimensionality and improve model performance.

9. Model Performance Metrics

9.1 Classification Metrics

- **Accuracy**: The proportion of predictions that are correct out of all predictions made.
- **Precision**: The proportion of true positives out of all predicted positives (true positives plus false positives).
- **Recall (Sensitivity)**: The proportion of true positives out of all actual positives (true positives plus false negatives).
- **F1 Score**: The harmonic mean of Precision and Recall, balancing both metrics.
- **Specificity**: The proportion of true negatives out of all actual negatives (true negatives plus false positives).

9.2 Regression Metrics

- **Mean Absolute Error (MAE)**: Average absolute difference between predictions and actual values.
- **Mean Squared Error (MSE)**: The average of the squared differences between the predicted and actual values. MSE heavily penalizes large errors and is commonly used for regression tasks.
- **R^2 (Coefficient of Determination)**: A measure of how much variance in the target variable is explained by the model. R^2 ranges from 0 to 1, where 1 indicates a perfect fit.

Bonus: What You Might Have Missed

Here are a few optional but relevant services to keep in mind:

- **Amazon Fraud Detector**: Detects potential fraudulent online activities using ML.
- **Amazon Lookout for Equipment**: Identifies abnormal equipment behavior to minimize downtime.
- **Amazon HealthLake**: Stores, transforms, queries, and analyzes health data at scale using ML.

Conclusion

Use this cheatsheet as a reference to rapidly review the essential knowledge about for the AWS Certified Machine Learning Engineer — Associate MLA-C01 exam. Happy Certification!

Written by Rafa Xu

7 followers · 3 following

Follow

AWS Cloud Architect with a passion for Cloud Computing, Software Development, and Data-driven solutions—now exploring the forefront of Generative AI.

Responses (1)